

# Case Study: Azure → Snowflake with Snowpark, then Power BI

Aarthi. P - 22/10/25

## Scenario:

You're the data engineer at ItTechGenie Retail. Sales teams drop monthly CSVs into an Azure Storage container.

## You must:

- upload the CSV to Azure
- ingest it into Snowflake using Snowpark
- model it into proper database/schema/table
- build a quick Power BI report for business users.

## Goal:

Ensure that business users have access to clean, structured, and visualized sales data for reporting and analysis.

## Step 1: Azure Storage Setup

### Create an Azure Storage Account:

- Go to the Azure Portal → Storage Accounts → Create.
- Choose the subscription, resource group, storage account name, and region.

The screenshot displays the Microsoft Azure portal interface for a storage account named 'storagesaccounts1'. The top navigation bar includes the Microsoft Azure logo, a search bar, and the user profile 'azuser4828\_mml.local@...'. The main content area shows the 'Overview' tab for the storage account, with a left-hand navigation pane listing various services like Activity log, Tags, Diagnose and solve problems, Access Control (IAM), Data migration, Events, Storage browser, Storage Mover, Partner solutions, Resource visualizer, Data storage, Containers, File shares, Queues, Tables, Security + networking, Data management, and Settings.

The 'Overview' section provides key details about the storage account:

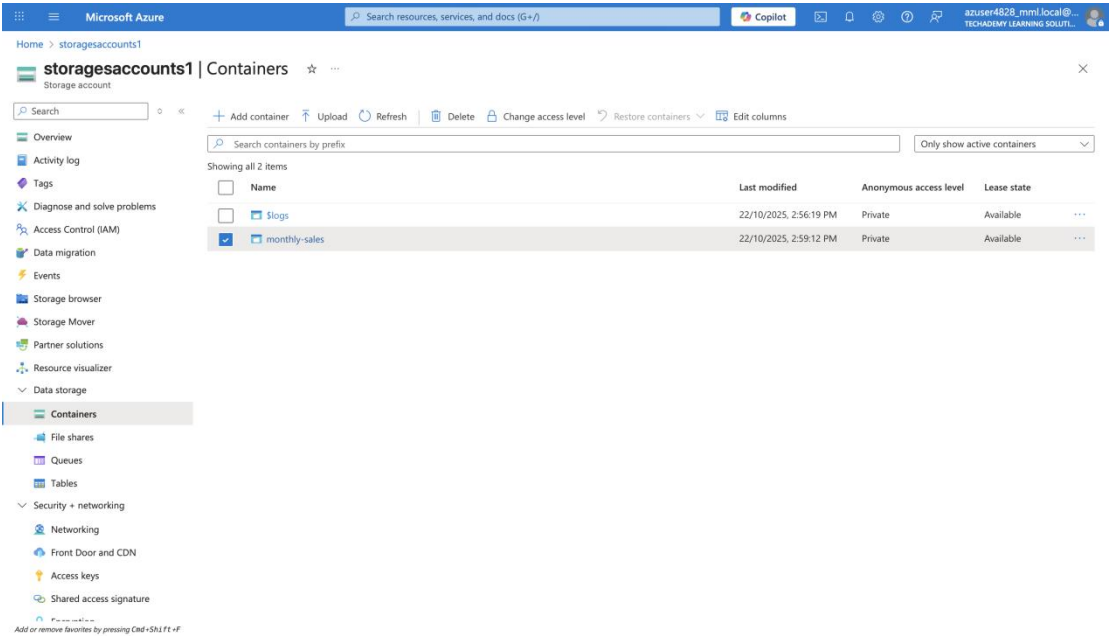
- Resource group: rg-azuser4828\_mml.local-WSHRo
- Location: southindia
- Subscription: MML\_Learners
- Subscription ID: 2a3c6418-97b9-4d96-a24b-2c2d7633d375
- Disk state: Available
- Tags: Add tags
- Performance: Standard
- Replication: Locally-redundant storage (LRS)
- Account kind: StorageV2 (general purpose v2)
- Provisioning state: Succeeded
- Created: 22/10/2025, 2:55:55 PM

The 'Properties' tab is active, showing configuration options for Blob service, Security, and Networking.

Property	Value
<strong>Blob service</strong>	
Hierarchical namespace	Disabled
Default access tier	Hot
Blob anonymous access	Disabled
Blob soft delete	Enabled (7 days)
Container soft delete	Enabled (7 days)
Versioning	Disabled
Change feed	Disabled
NFS v3	Disabled
Allow cross-tenant replication	Disabled
Storage tasks assignments	None
<strong>Security</strong>	
Require secure transfer for REST API operations	Enabled
Storage account key access	Enabled
Minimum TLS version	Version 1.2
Infrastructure encryption	Disabled
<strong>Networking</strong>	
Public network access	Enabled
Public network access scope	Enable from all networks
Private endpoint connections	0
Network routing	Microsoft network routing
Endpoint type	Standard

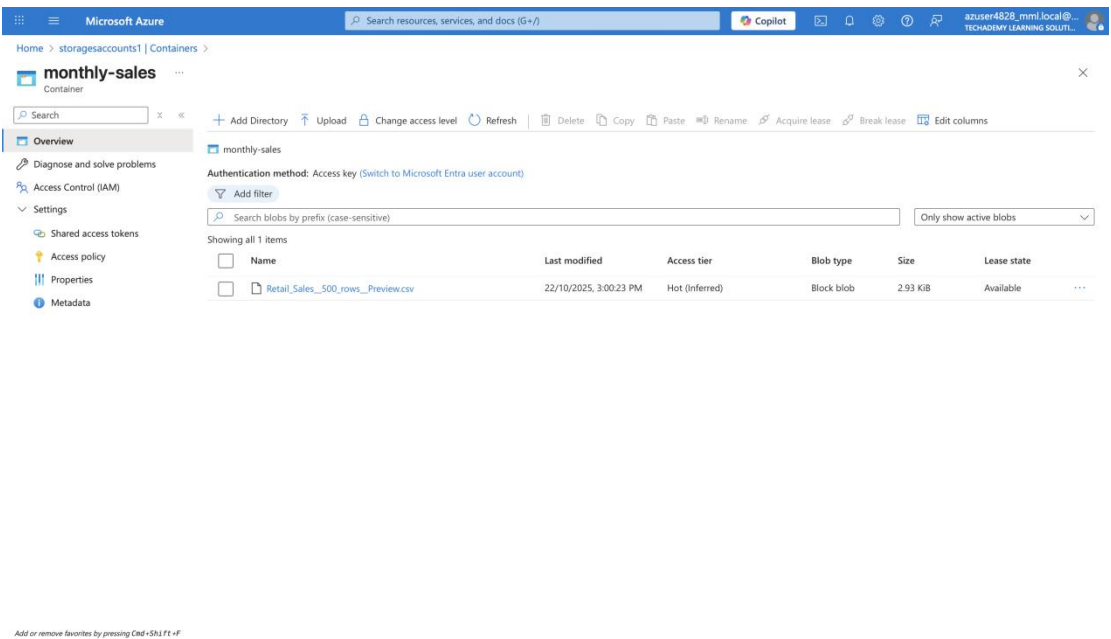
Create a Container:

- Inside the storage account, create a container (e.g., sales-monthly).
- Set access level to private.



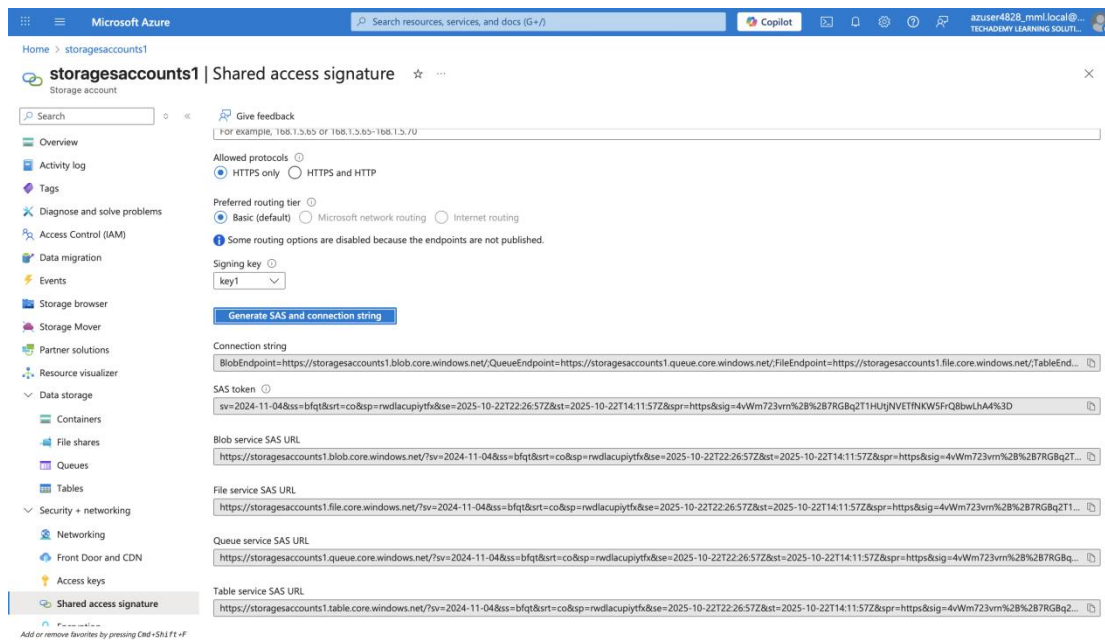
Upload CSV Files:

- Upload monthly CSVs (e.g., Retail\_Sales\_\_500\_rows\_\_Preview.csv).



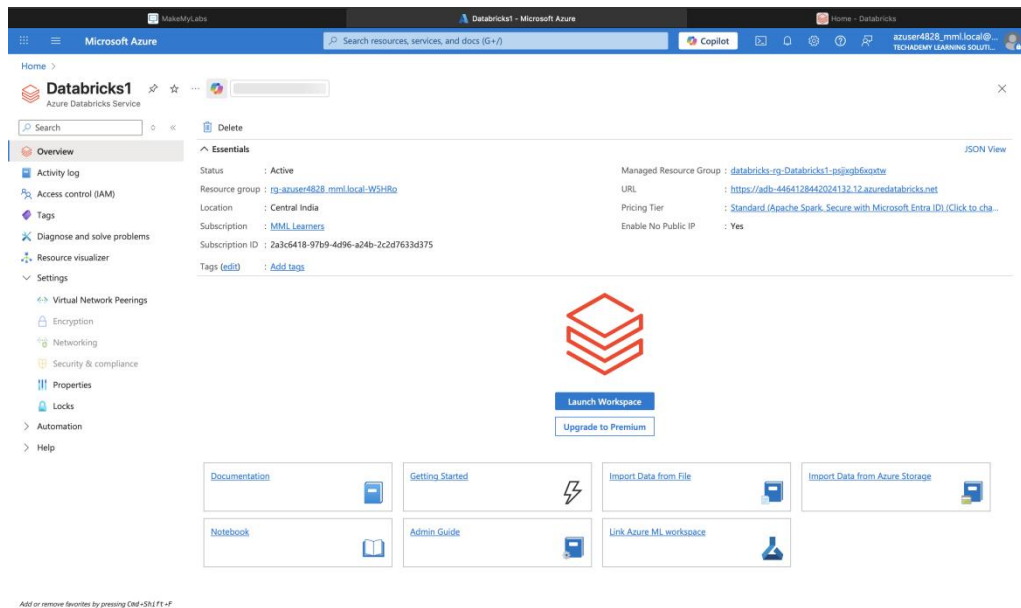
## Generate SAS Token for Access:

- Click on the file or container → Generate SAS.
- Set permissions (Read, List, Write if needed).
- Set expiry date/time.
- Copy the full URL including SAS token for later use.

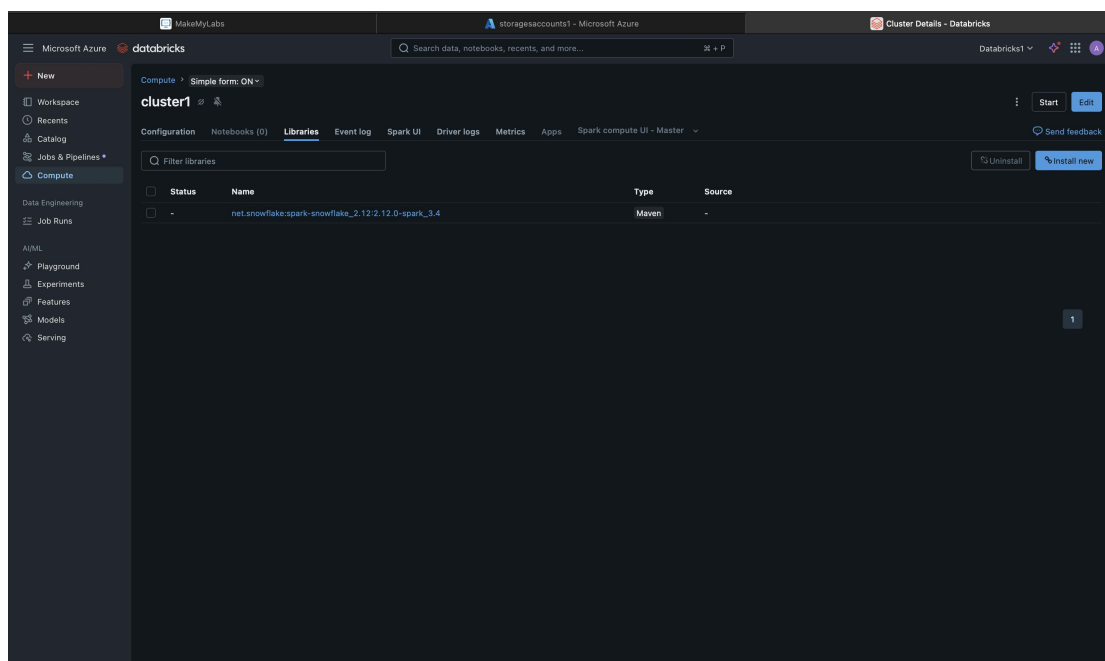


## Step 2: Azure Databricks Setup:

- Create a Databricks Workspace (if not already available).



- Launch a Cluster with appropriate size.



- Install Required Libraries.

Step 2: External stage creation:

```
CREATE OR REPLACE DATABASE SNOWPIPE;
CREATE OR REPLACE SCHEMA RETAIL_SALES;
```

```
CREATE OR REPLACE FILE FORMAT azure_csv_format
TYPE = 'CSV'
FIELD_OPTIONALLY_ENCLOSED_BY = ''
SKIP_HEADER = 1;
```

```
CREATE OR REPLACE STAGE azure_sales_stage
```

```

URL = 'azure://storageaccounts1.blob.core.windows.net/monthly-sales'
CREDENTIALS = (AZURE_SAS_TOKEN = 'sv=2024-11-04&ss=bfqt&srt=sco&sp=rwdlacupiytfx&se=2025-10-22T18:20:22Z&st=2025-10-22T10:05:22Z&spr=https&sig=%2FZqM2TSKB42aMjS5nnuMUnQ4ZIXAAMjFbdecscIJ988%3D')

FILE_FORMAT = azure_csv_format;

```

### Step 3.1: Verification:

```

LIST @azure_sales_stage;

SHOW STAGES LIKE 'azure_sales_stage';

```

The screenshot shows the Snowflake SQL Editor interface. The left sidebar displays the database structure, with 'SALES\_ORDERS' selected under the 'RETAIL\_SALES' schema. The main editor shows the following SQL script:

```

1 CREATE OR REPLACE DATABASE SNOWPIPE;
2 CREATE OR REPLACE SCHEMA RETAIL_SALES;
3
4 CREATE OR REPLACE FILE FORMAT azure_csv_format
5 TYPE = 'CSV'
6 FIELD_OPTIONALLY_ENCLOSED_BY = ''
7 SKIP_HEADER = 1;
8
9 CREATE OR REPLACE STAGE azure_sales_stage
10 URL = 'azure://storageaccounts1.blob.core.windows.net/monthly-sales'
11 CREDENTIALS = (AZURE_SAS_TOKEN = 'sv=2024-11-04&ss=bfqt&srt=sco&sp=rwdlacupiytfx&se=2025-10-22T18:20:22Z&st=2025-10-22T10:05:22Z&spr=https&sig=%2FZqM2TSKB42aMjS5nnuMUnQ4ZIXAAMjFbdecscIJ988%3D')
12 FILE_FORMAT = azure_csv_format;
13
14 LIST @azure_sales_stage;
15
16 SHOW STAGES LIKE 'azure_sales_stage';

```

The bottom pane shows the results of the 'LIST @azure\_sales\_stage;' command, indicating a successful stage creation:

name	size	md5	last_modified
azure://storageaccounts1.blob.core.windows.net/monthly-sales/Retail_Sales_500_rows_Preview.csv	3004	a38163e0b0e75a2fc23a050902	Wed, 22 Oct 2025 09:30:23 GMT

### Step 4: Table creation in snowlake:

```

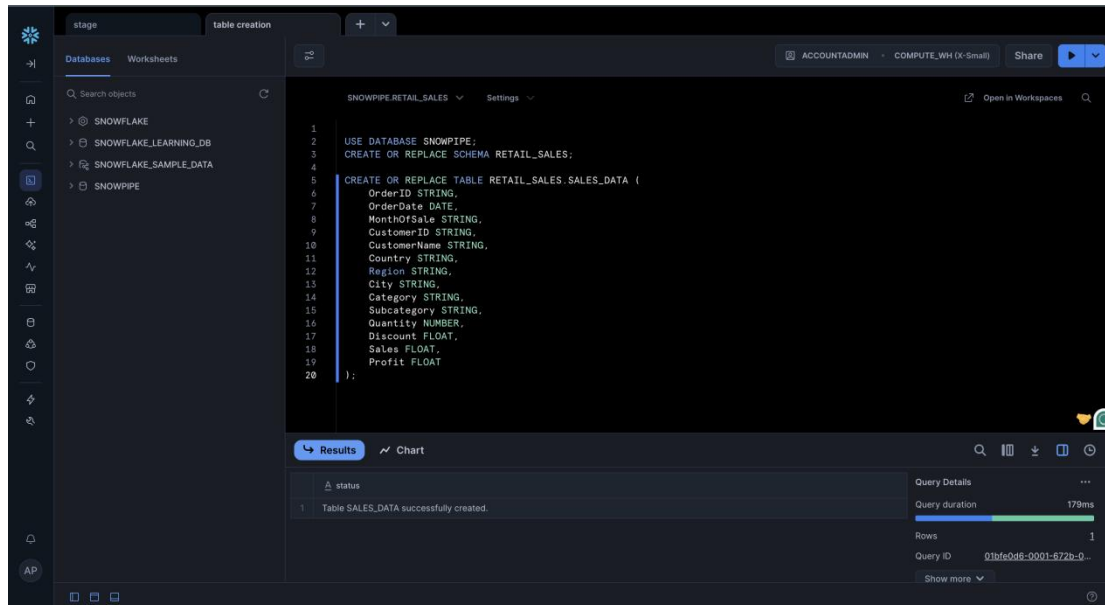
USE DATABASE SNOWPIPE;
USE SCHEMA RETAIL_SALES;

CREATE OR REPLACE TABLE
RETAIL_SALES.SALES_ORDERS (
  OrderID STRING,
  OrderDate DATE,
  MonthOfSale STRING,
  CustomerID STRING,
  CustomerName STRING,
  Country STRING,
  Region STRING,
  City STRING,
  Category STRING,
  Subcategory STRING,
  Quantity NUMBER,
  Discount FLOAT,
  Sales FLOAT,
  Profit FLOAT);

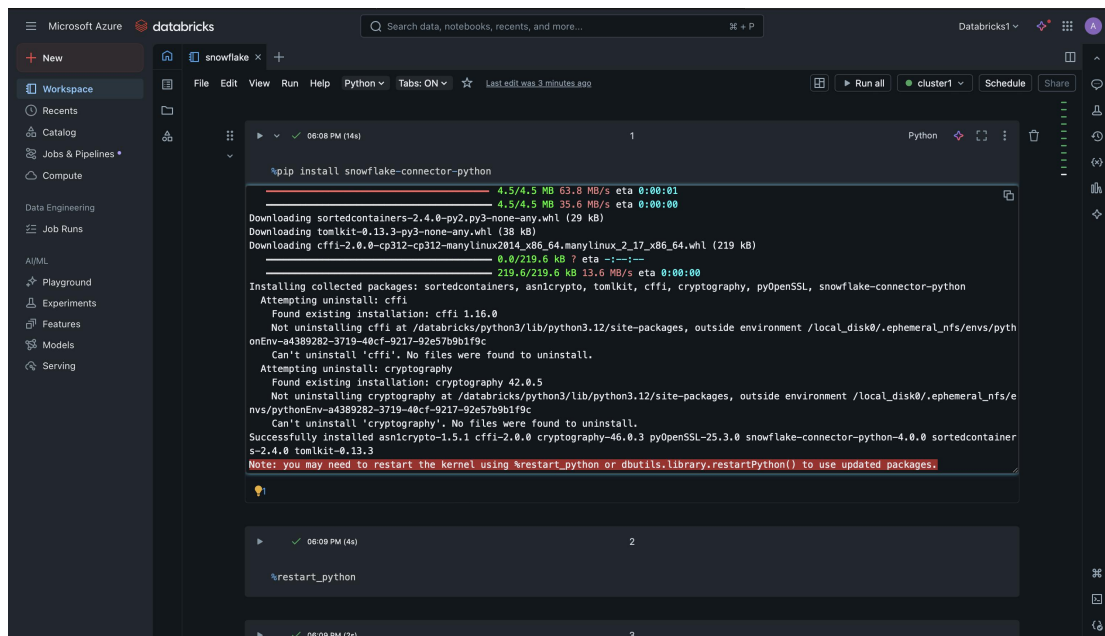
```

## Step 4.1: Data Loading:

```
COPY INTO RETAIL_SALES.SALES_ORDERS
FROM @azure_sales_stage
FILE_FORMAT = (FORMAT_NAME = azure_csv_format)
ON_ERROR = 'CONTINUE';
```



## Step 5: Snowpark - Data Transformation:



## Step 6: Connect Storage account with databricks:

The screenshot shows a Databricks workspace with a notebook titled 'snowflake'. The notebook contains the following code and output:

```
import pandas as pd
from snowflake.connector import connect
from snowflake.connector.pandas_tools import write_pandas
```

```
url = "https://storageaccounts1.blob.core.windows.net/monthly-sales/Retail_Sales_500_rows_Preview.csv?sv=2024-11-04&ss=bfgt&srt=co&sp=rwdlacupiytfx&se=2025-10-22T18:20:22Z&st=2025-10-22T18:05:22Z&spr=https&sig=%2FzqM2TSK842mJS5nnuWu04ZLXAAMjFbdecscI3988%30"
```

```
df = pd.read_csv(url)
```

```
df: pandas.core.frame.DataFrame = [OrderID: object, OrderDate: object ... 12 more fields]
```

```
df.head(5)
```

	OrderID	OrderDate	MonthOfSale	CustomerID	CustomerName	Country	Region	City	Category	Subcategory	Quantity	Discount	Sales	Profit
0	ORD-SF8D6FOC	2024-10-08	2024-10	CUST1000	Ananya Sharma	India	South	Mumbai	Office Supplies	Paper	9	0.00	2700.0	780.43
1	ORD-BF0078E4	2024-08-11	2024-08	CUST1001	Aarav Iyer	India	Central	Lucknow	Technology	Networking	4	0.15	27200.0	4135.60
2	ORD-86CD98A3	2024-06-15	2024-06	CUST1002	Arjun Sharma	USA	East	Kolkata	Furniture	Tables	4	0.10	31500.0	5676.96
3	ORD-FB0CD0D9	2024-12-18	2024-12	CUST1003	Ananya Das	India	North	Kolkata	Office Supplies	Appliances	9	0.00	38000.0	11783.22
4	ORD-EF35596B	2024-10-27	2024-10	CUST1004	Ishaan Bhat	UK	Central	Chennai	Furniture	Storage	4	0.00	24000.0	4189.98

Step 7: Snowflake connection:

The screenshot shows a Databricks workspace with a notebook titled 'snowflake'. The notebook contains the following code and output:



```
conn = connect(
    user="AARTHI",
    password="Snowflake@9104",
    account="XFCJLB0-UH96778",
    warehouse="COMPUTE_WH",
    database="SNOWPIPE",
    schema="RETAIL_SALES"
)
```

```
result = write_pandas(
    conn=conn,
    df=df,
    table_name="SALES_ORDERS",
    schema="RETAIL_SALES",
    database="SNOWPIPE",
    auto_create_table=True,
    quote_identifiers=False
)
```

```
result[0]
```

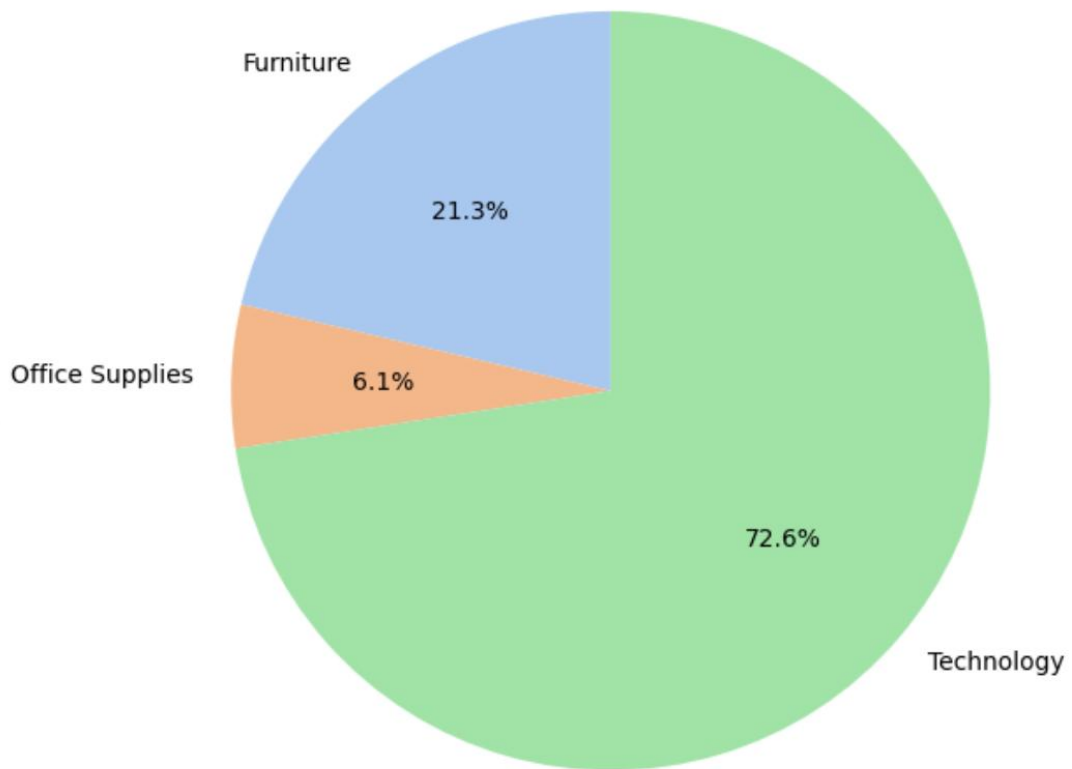
```
True
```

Step 8: Visualization:





Sales Distribution by Category



Total Sales by Region

